

HR Analytics Project- Understanding the Attrition in HR



Problem Statement

Every year a lot of companies hire a number of employees. The companies invest time and money in training those employees, not just this but there are training programs within the companies for their existing employees as well.

The objective of the model to increase the effectiveness of their employees and reduce the time and money investing in employees.

HR Analytics:

Human resource analytics (HR analytics) is an area in the field of analytics that refers to applying analytic processes to the human resource department of an organization in the hope of improving employee performance and therefore getting a better return on investment. HR analytics does not just deal with gathering data on employee efficiency. Instead, it aims to provide insight into each process by gathering data and then using it to make relevant decisions about how to improve these processes.

Attrition in HR

Attrition in human resources refers to the gradual loss of employees overtime. In general, relatively high attrition is problematic for companies. HR professionals often assume a leadership role in designing company compensation programs, work culture, and motivation systems that help the organization retain top employees

How does Attrition affect companies? and how does HR Analytics help in analysing attrition? We will discuss the first question here and for the second question, we will write the code and try to understand the process step by step.

Attrition affecting Companies

A major problem in high employee attrition is its cost to an organization. Job postings, hiring processes, paperwork, and new hire training are some of the common expenses of losing employees and replacing them. Additionally, regular employee turnover prohibits your organization from increasing its collective knowledge base and experience over time. This is especially concerning if your business is customer-facing, as customers often prefer to interact with familiar people. Errors and issues are more likely if you constantly have new workers.

Importing the Libraries

```
import pandas as pd
import numpy as np
import seaborn as sns
from scipy import stats
import warnings
warnings.filterwarnings('ignore')
from matplotlib import pyplot as plt
from scipy.stats import zscore
#data preprocessing
from sklearn import preprocessing
from sklearn.preprocessing import MinMaxScaler
#Over Sampling the data using SMOTE
from imblearn.over_sampling import SMOTE
#modelling
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import cross_val_score
from matplotlib import pyplot
from sklearn.svm import SVC
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.tree import DecisionTreeClassifier
```

EXPLORATORY DATA ANALYSIS

In [3]:	1 df.describe()
Out[3]:	
	Education EmployeeCount EmployeeNumber EnvironmentSatisfaction HourlyRate JobInvolvement JobLevel ... RelationshipSatisfaction StandardHours
	1470.000000 1470.0 1470.000000 1470.000000 1470.000000 1470.000000 1470.000000 ... 1470.000000 1470.0
	2.912925 1.0 1024.865306 2.721769 65.891156 2.729932 2.063946 ... 2.712245 80.0
	1.024165 0.0 602.024335 1.093082 20.329428 0.711561 1.106940 ... 1.081209 0.0
	1.000000 1.0 1.000000 1.000000 30.000000 1.000000 1.000000 ... 1.000000 80.0
	2.000000 1.0 491.250000 2.000000 48.000000 2.000000 1.000000 ... 2.000000 80.0
	3.000000 1.0 1020.500000 3.000000 66.000000 3.000000 2.000000 ... 3.000000 80.0
	4.000000 1.0 1555.750000 4.000000 83.750000 3.000000 3.000000 ... 4.000000 80.0
	5.000000 1.0 2068.000000 4.000000 100.000000 4.000000 5.000000 ... 4.000000 80.0

Employee's average number of years at company is 7.

Mean is not equal to median stating that the data is not normally distributed. Most normally distributes column is Daily rate where mean is almost equal to median

Out[4]:	Age	int64
	Attrition	object
	BusinessTravel	object
	DailyRate	int64
	Department	object
	DistanceFromHome	int64
	Education	int64
	EducationField	object
	EmployeeCount	int64
	EmployeeNumber	int64
	EnvironmentSatisfaction	int64
	Gender	object
	HourlyRate	int64
	JobInvolvement	int64
	JobLevel	int64
	JobRole	object
	JobSatisfaction	int64
	MaritalStatus	object
	MonthlyIncome	int64
	MonthlyRate	int64
	NumCompaniesWorked	int64
	Over18	object
	OverTime	object
	PercentSalaryHike	int64
	PerformanceRating	int64
	RelationshipSatisfaction	int64
	StandardHours	int64
	StockOptionLevel	int64
	TotalWorkingYears	int64
	TrainingTimesLastYear	int64
	WorkLifeBalance	int64
	YearsAtCompany	int64

Numeric variables:

- Related to personal information: age, distance_from_home, employee_number
- Related to income: hourly_rate, daily_rate, monthly_rate, monthly_income, percent_salary_hike

Related to duration in company: years_at_company, years_in_current_role, years_since_last_promotion, years_with_curr_manager, total_working_years

num_companies_worked, standard_hourstraining_times_last_year, employee_count

Categorical variables:

- Binary variables: attrition(target variable), gender, over18, over_time
- Nominal variables: department, education_field, job_role, marital_status
- Ordinal variables:

- Ordinal regarding satisfaction and performance: environment_satisfaction, job_satisfaction, relationship_satisfaction, work_life_balance, job_involvement, performance_rating
- Other ordinal: business_travel, education, job_level, stock_option_level

```

Handling null values

In [5]: 1 df.isnull().sum()

Out[5]: Age 0
Attrition 0
BusinessTravel 0
DailyRate 0
Department 0
DistanceFromHome 0
Education 0
EducationField 0
EmployeeCount 0
EmployeeNumber 0
EnvironmentSatisfaction 0
Gender 0
HourlyRate 0
JobInvolvement 0
JobLevel 0
JobRole 0
JobSatisfaction 0
MaritalStatus 0
MonthlyIncome 0
MonthlyRate 0
NumCompaniesWorked 0
Over18 0
OverTime 0
PercentSalaryHike 0
PerformanceRating 0
RelationshipSatisfaction 0
StandardHours 0

```

This dataset has no null values

```

In [24]: 1 for col in df:
2         print(col)
3         print(df[col].value_counts())
4         print()

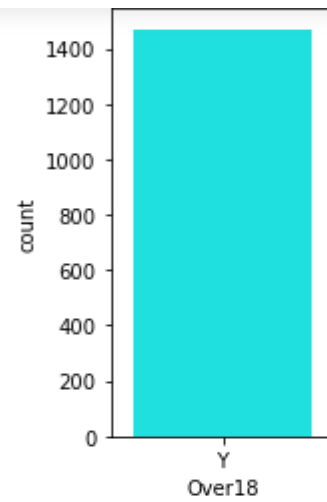
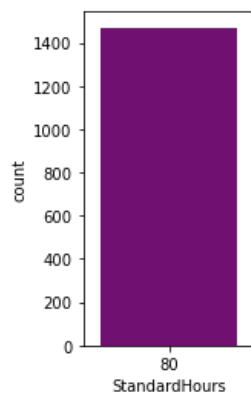
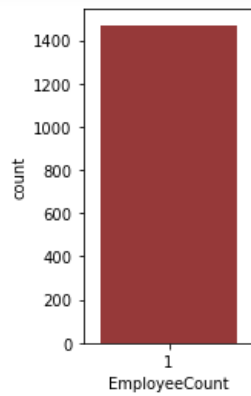
1 1470
Name: EmployeeCount, dtype: int64

EmployeeNumber
2048 1
1368 1
1364 1
1363 1
1362 1
..
648 1
647 1
645 1
644 1
2046 1
Name: EmployeeNumber, Length: 1470, dtype: int64

```

Displaying value count of unique value in each feature. To identify column having single unique value

UNI VARIATE ANALYSIS



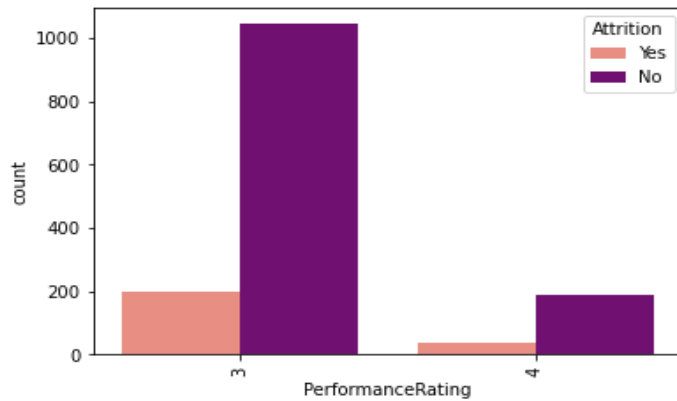
The features standard hours, over18 and employee count has only one value so it won't create any impact on the target feature Attrition.

All the three columns having single value so I'm going to dropping it from the given dataset

BI VARIATE ANALYSIS (Categorical columns vs Target)

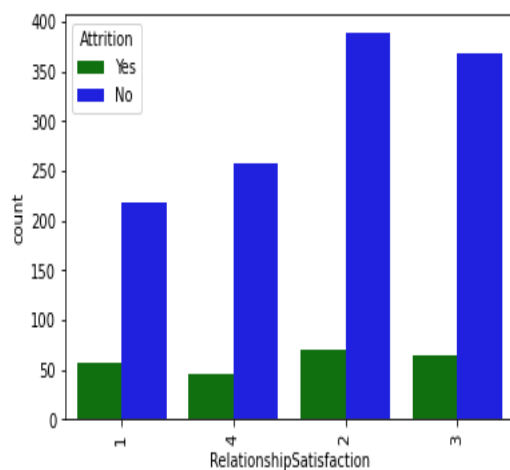
```
1 l = list(df['PerformanceRating'].unique())
2 col_1=["salmon","purple"]
3 chart = sns.countplot(df["PerformanceRating"],palette=col_1,hue=df.Attrition)
4 chart.set_xticklabels(labels=l, rotation=90)
```

```
[Text(0, 0, '3'), Text(1, 0, '4')]
```

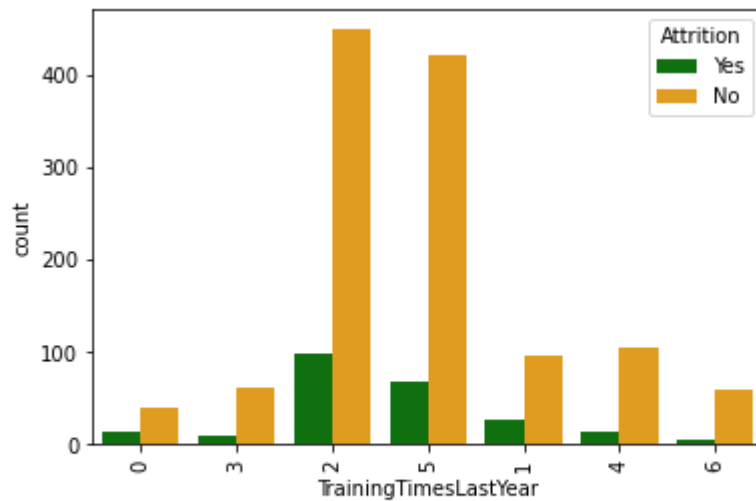


```
1 the employees who got performance rating as 3 are having less attrition rate.
```

```
: [Text(0, 0, '1'), Text(1, 0, '4'), Text(2, 0, '2'), Text(3, 0, '3')]
```

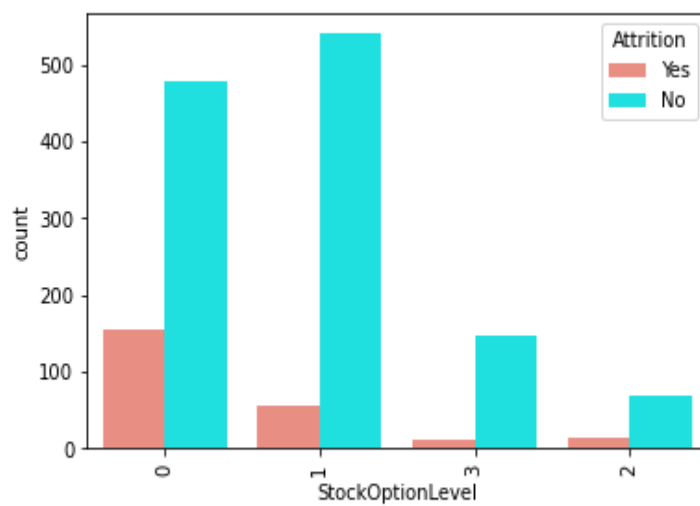


from the above chart its apparent that the employee who having high relationship satisfaction and low relationship satisfaction are having low Attrition rate.so we cannot predict the target coloumn with this value

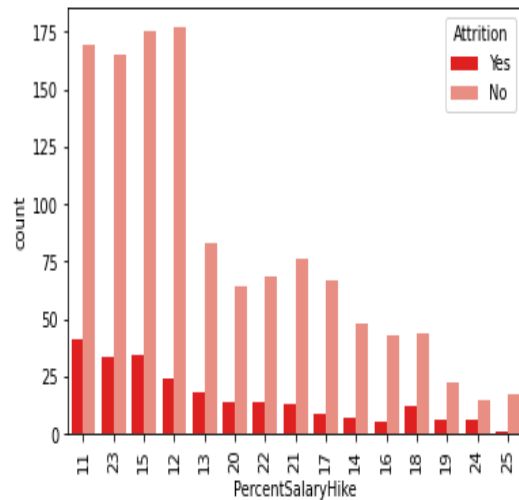


Employee who are trained 2 and 5 times a year are having less attrition rate

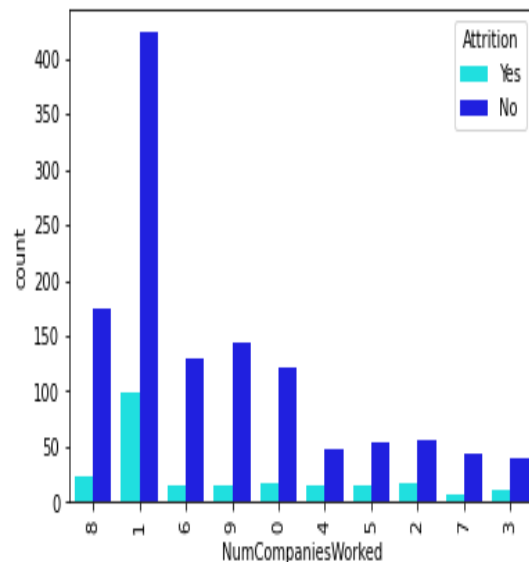
```
]: [Text(0, 0, '0'), Text(1, 0, '1'), Text(2, 0, '3'), Text(3, 0, '2')]
```



the employees who having less stocks are having low attrition rate.



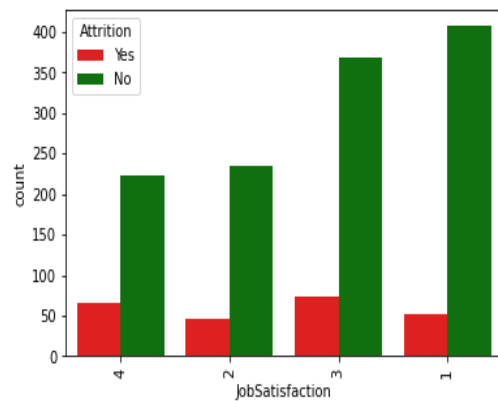
the employees who all are got hike 11 to 15% are having less attrition rate.the employees who got hike between 18 to 20% having high attrition rate .we can predict the target column using this feature



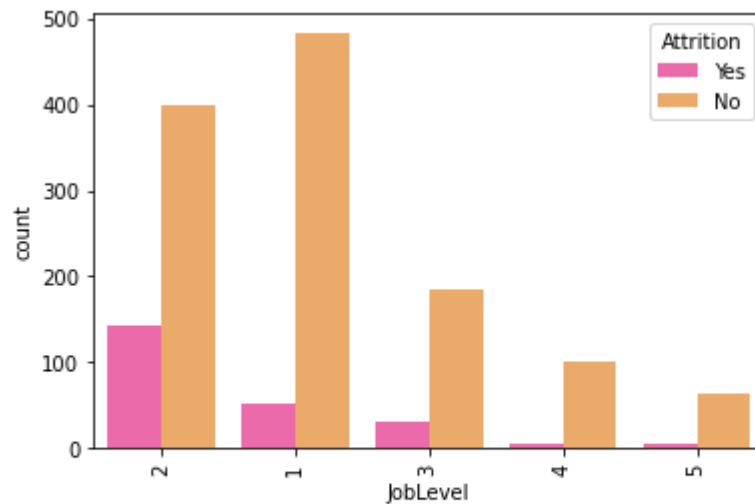
from the above chart its apparent that both the employee who having high satisfaction and low satisfaction are having low Attrition rate.so we cannot predict the target coloumn with this value

from the above chart its apparent that the employees who worked in only one company are having low Attrition rate.

[Text(0, 0, '4'), Text(1, 0, '2'), Text(2, 0, '3'), Text(3, 0, '1')]



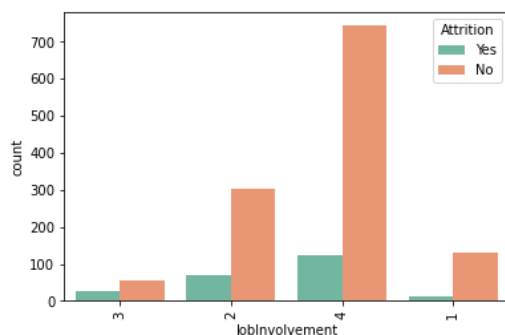
from the above chart its apparent that the employee who having high job satisfaction and low job satisfaction are having low Attrition rate.so we cannot predict the target column with this value



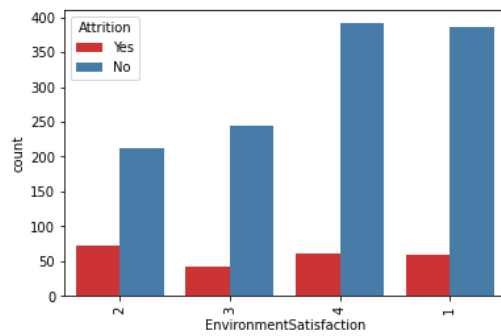
from the above chart its apparent that the entry level employees are having low Attrition rate.

```
33]: 1 = list(df['JobInvolvement'].unique())
      2 chart = sns.countplot(df["JobInvolvement"],palette="Set2",hue=df.Attrition)
      3 chart.set_xticklabels(labels=1, rotation=90)
```

33]: [Text(0, 0, '3'), Text(1, 0, '2'), Text(2, 0, '4'), Text(3, 0, '1')]

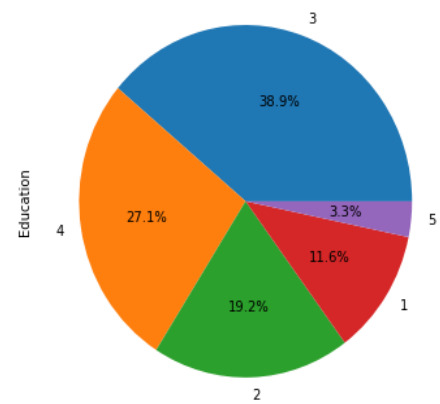
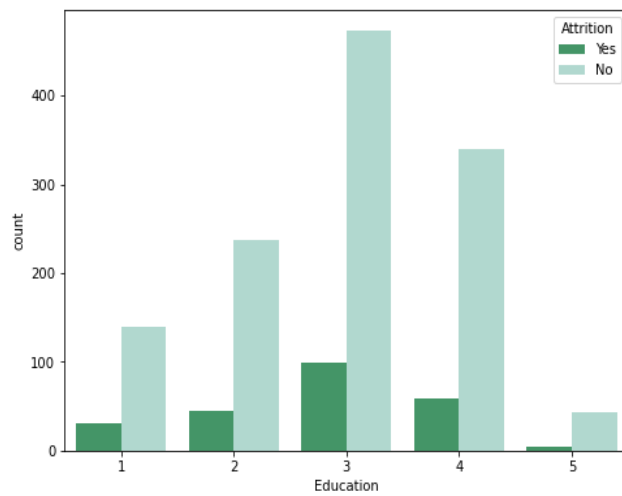


from the above chart its apparent that the employee who all are highly involved in the job having low Attrition rate

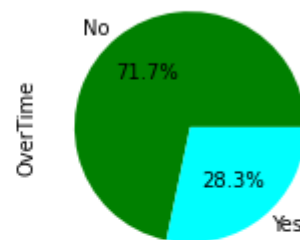
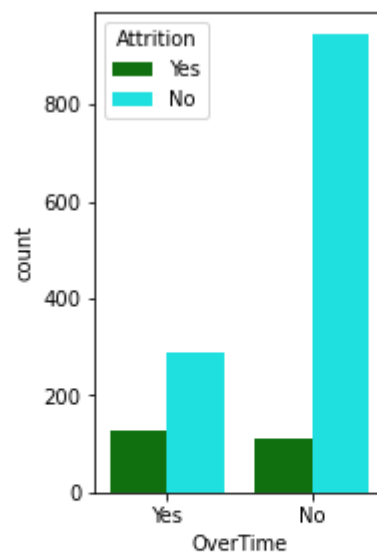


from the above chart its apparent that both the employee who having high satisfaction and low satisfaction are having low Attrition rate.so we cannot predict the target column with this value

```
<AxesSubplot:ylabel='Education'>
```



The employees of 38.9% belongs to education category 3.compare to other category 3 has less percentage of people moving out of the company



Employees who all are not working overtime has low attrition rate

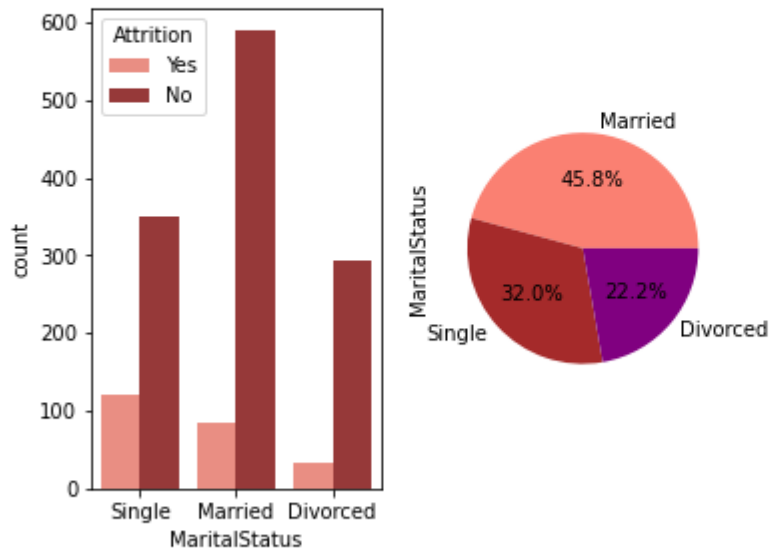
28.3% employees are willing to work in overtime

```

Married    673
Single    470
Divorced   327
Name: MaritalStatus, dtype: int64

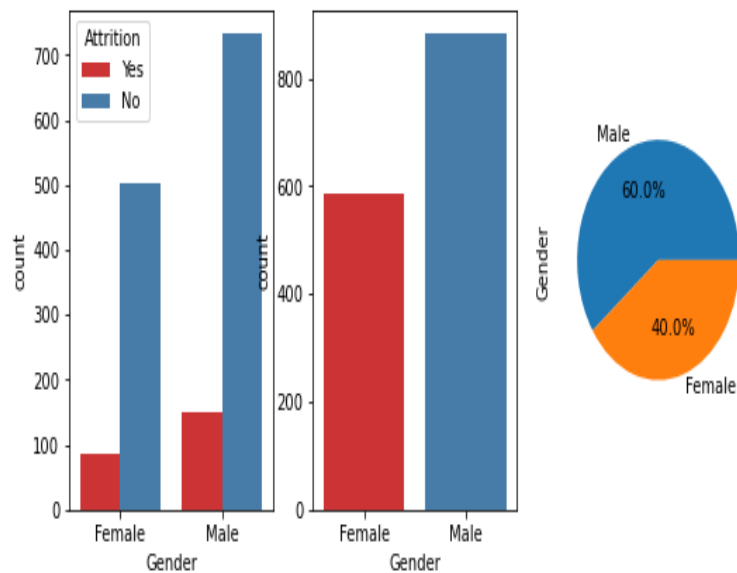
```

```
ut[12]: <AxesSubplot:ylabel='MaritalStatus'>
```

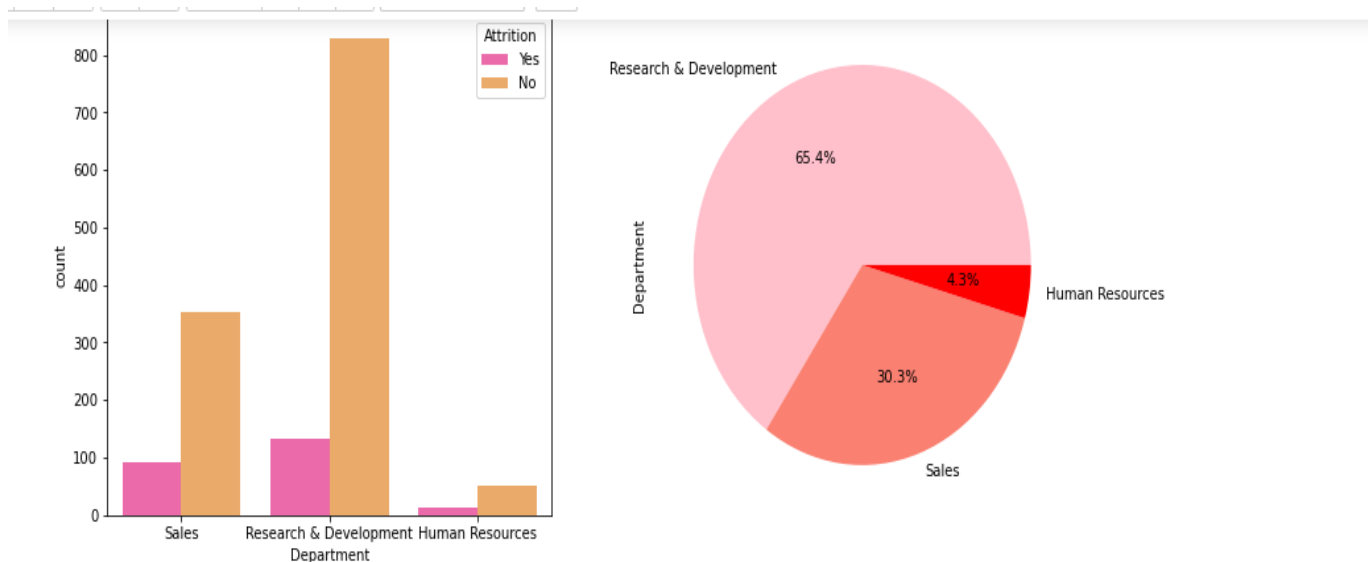


from this above chart its apparent that employees who all are married are having less attrition rate

```
]: <AxesSubplot:ylabel='Gender'>
```



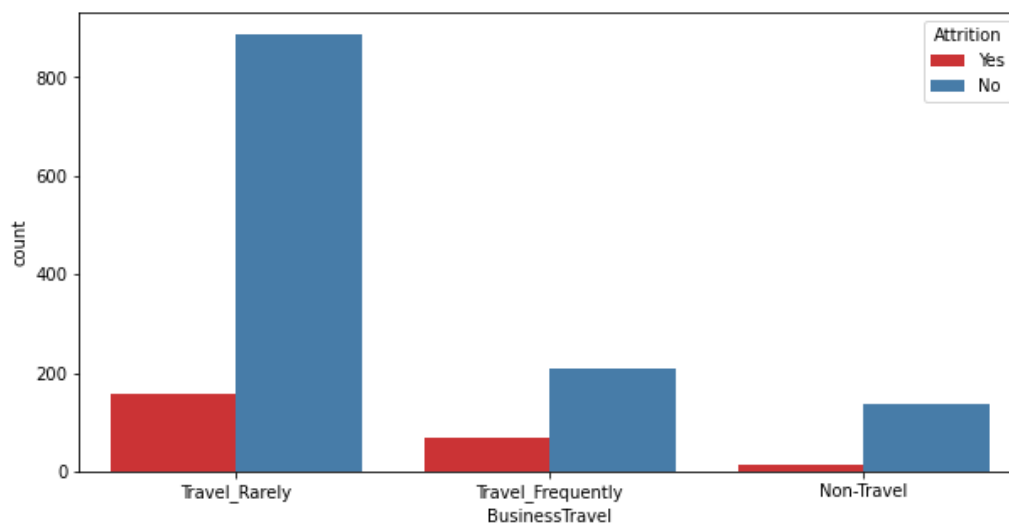
comparing the percentage of attrition out of 588 female only 88 people are quitting it means $\left(\frac{588-88}{588} \times 100\right)$ 15% female are leaving but in male out of 882 people more than 150 people are quitting it means $\left(\frac{882-150}{882} \times 100\right)$ 18% male are leaving



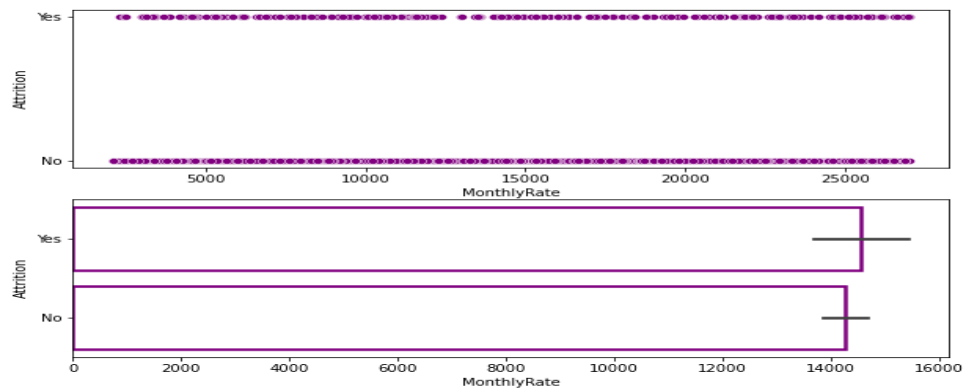
from the above plot its apparent that comparatively the employees who belongs to Research and Development will like to continue their job. The majority percentage(65.4%) of employee belongs to R&D department

```
[38]: 1 plt.figure(figsize=(10,5))
      2 sns.countplot(df.BusinessTravel,palette="Set1",hue=df.Attrition)
```

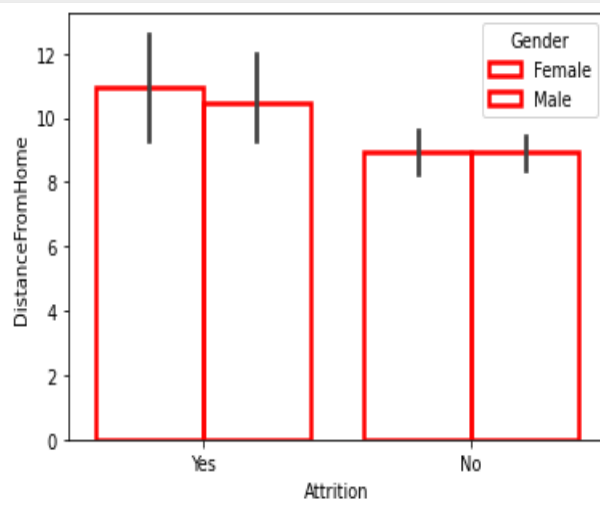
```
[38]: <AxesSubplot:xlabel='BusinessTravel', ylabel='count'>
```



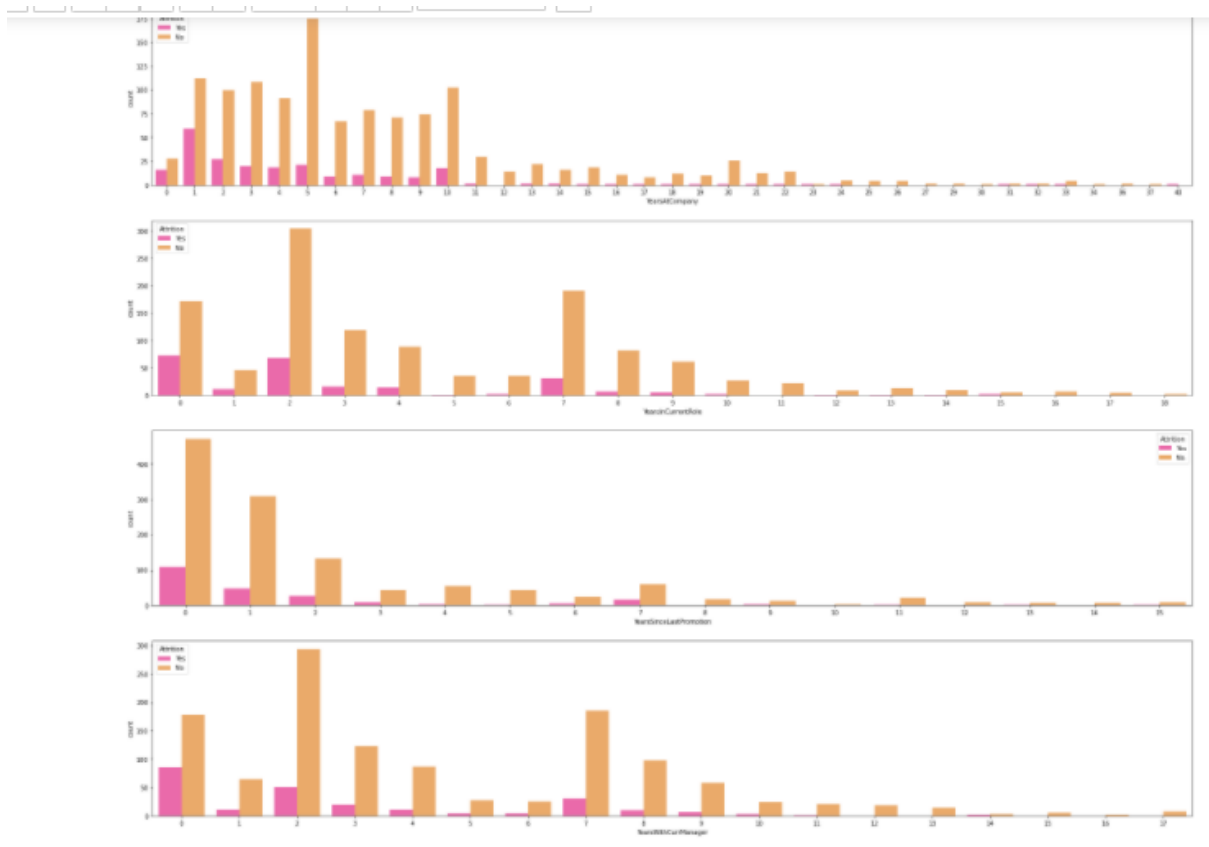
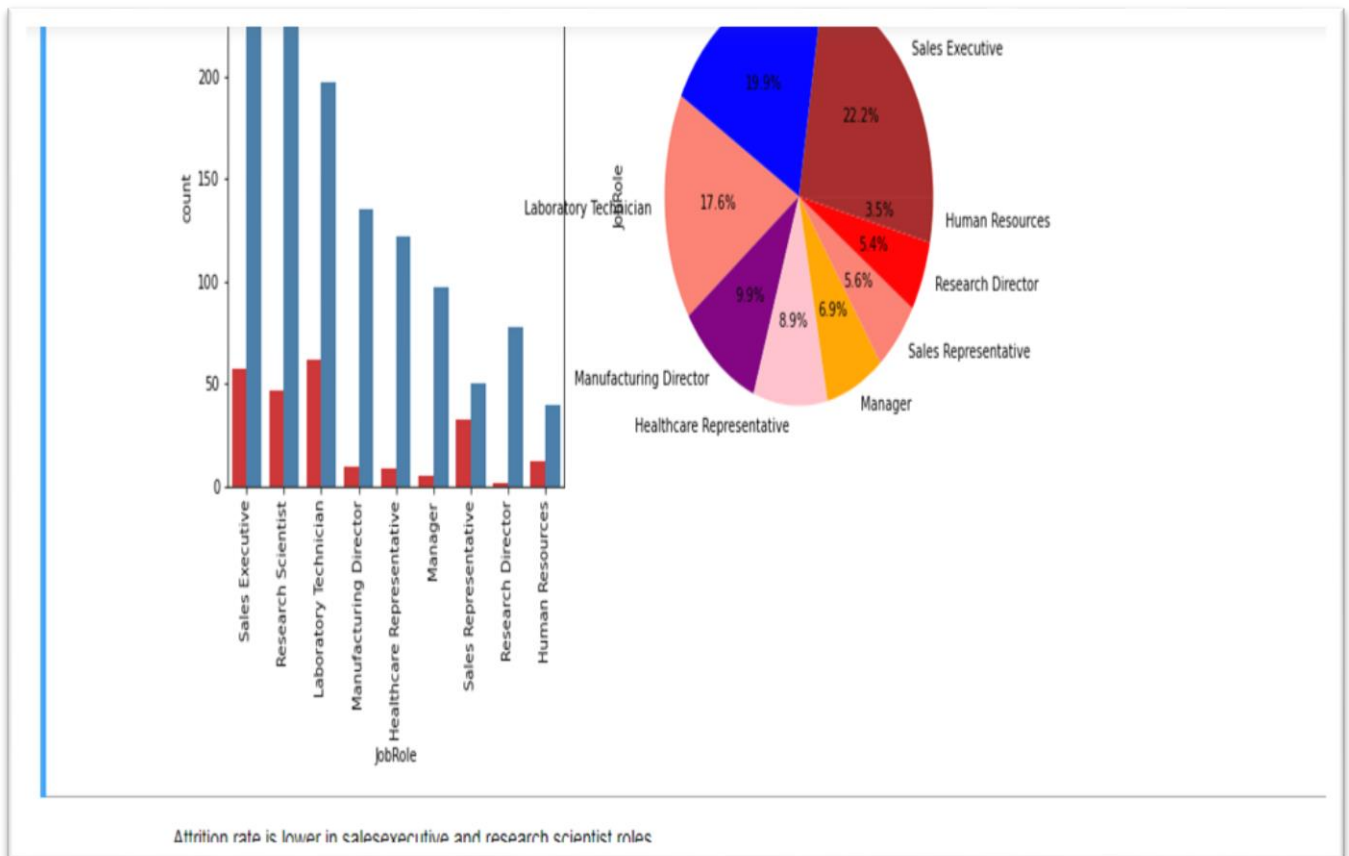
from the above plot its apparent that comparatively the employees who travel rarely will not resign their job



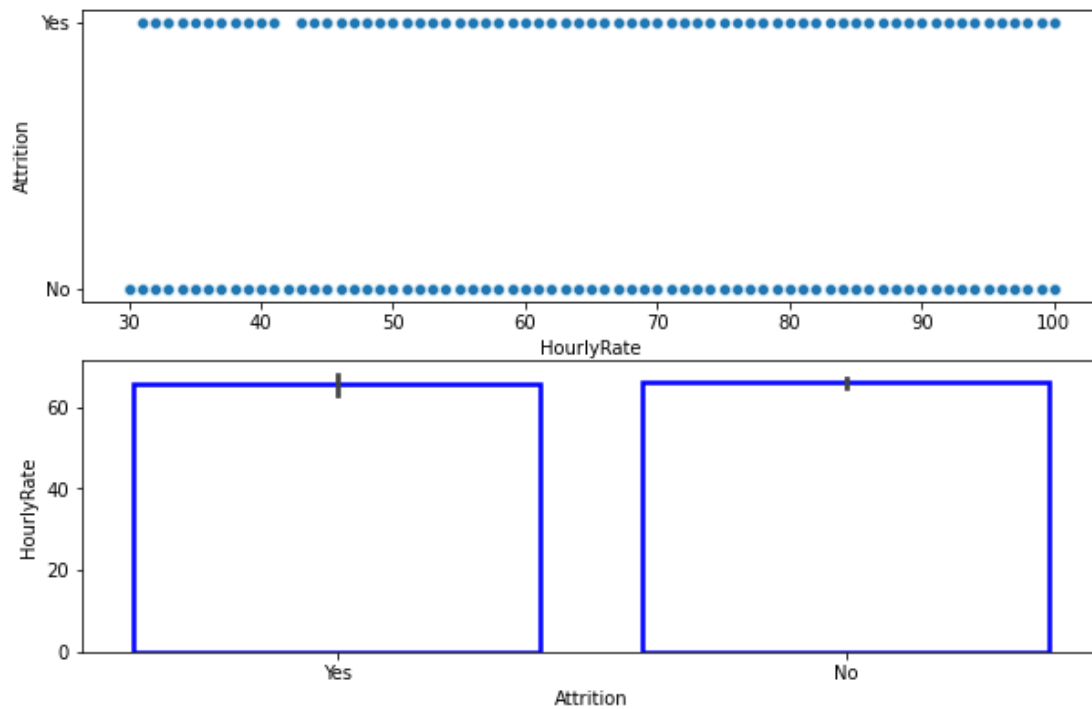
Employees having low monthlyrate are mostly resigning their job



Distance from home is not an important feature to create impact on Attrition feature. Distance from home is not impact on gender



```
J: <Axes>subplot(xlabel= Attrition , ylabel= HourlyRate >
```



HourlyRate is not an important feature to create impact on Attrition feature

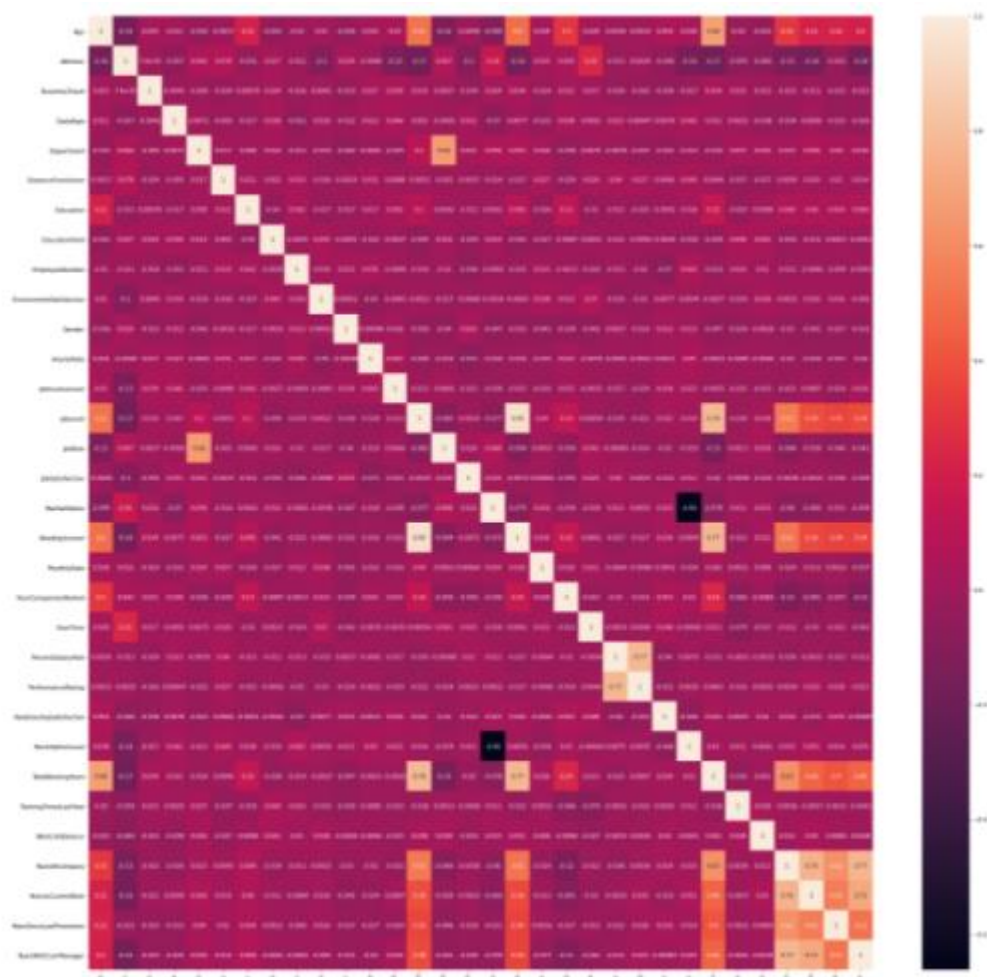
EDA CONCLUSION:

Employees who belongs to below category having less attritionon rate

- travel rarely
- who belongs to R&D department
- who belongs to life science and medical field
- female
- working as sales executive and research scientists
- unmarried
- not working over time
- moderate work life balance
- high job involvement
- working in single company
- performance rating:3
- employees who having 0 stocks
- low monthly income.

we cannot predict using relationship satisfaction,job satisfaction,Environment satisfaction features

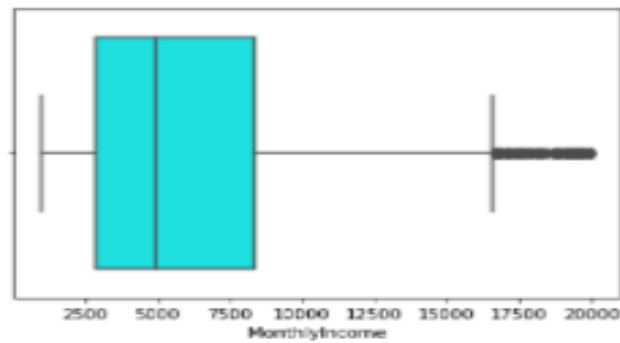
Correlation



Over time feature is highly correlated with attrition

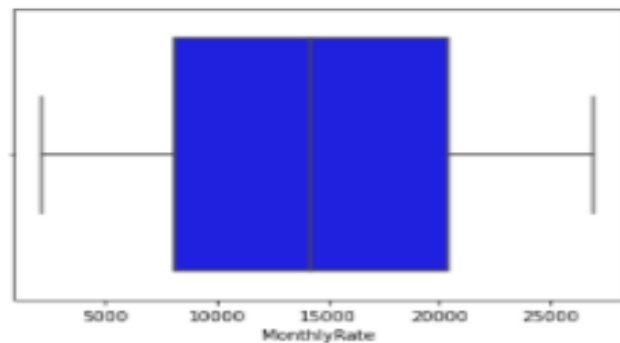
Handling outliers in numerical column:

```
Out[7]: <AxesSubplot:xlabel='MonthlyIncome'>
```



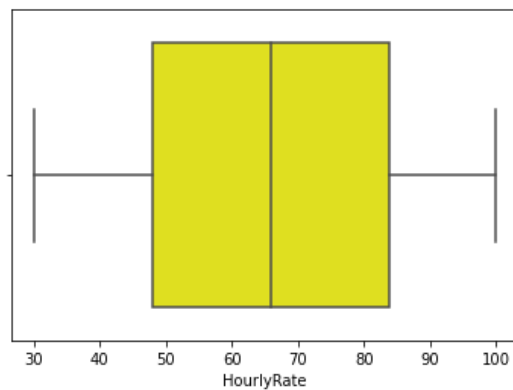
```
In [8]: 1 sns.boxplot(df['MonthlyRate'],color="blue")
```

```
Out[8]: <AxesSubplot:xlabel='MonthlyRate'>
```



```
In [9]: 1 sns.boxplot(df['DailyRate'],color="green")
```

```
Out[10]: <AxesSubplot:xlabel='HourlyRate'>
```



```
In [11]: 1 z1 = np.abs(stats.zscore(df_new['MonthlyIncome']))
2 print(z1)
```

```
[0.10834951 0.29171859 0.93765369 ... 0.07669019 0.23647414 0.44597809]
```

```
In [14]: 1 df_new['MonthlyIncome'] = df_new.MonthlyIncome[(z1<3)]
2 df_new.shape
```

```
Out[14]: (1470, 32)
```

_outliers are removed from numerical data monthly income

Data pre-processing:

The features standard hours, over18 and employee count has only single value so it won't create any impact on the target feature Attrition.

Employee Number feature is just an identifier and it's not required for modelling either. So I'm dropping these features

The features standardhours,over18 and employeecount has only one value so it wont create any impact on the target feature Attrition.

```
In [4]: 1 cols=['StandardHours','Over18','EmployeeCount']
        2 df_new=df.drop(cols,axis=1)
```

All the three columns having single value so I'm dropping it from the given dataset

```
In [4]: 1 df_new.shape
Out[4]: (1470, 32)
```

Encoding all categorical column into numerical column using label encoding technique

```
In [22]: 1 data_clean=df_new
        2 col_encoded=['Attrition','BusinessTravel','Department','EducationField','Gender','JobRole','MaritalStatus','OverTime']
```

```
In [23]: 1 from sklearn import preprocessing
        2 for col in col_encoded:
        3     label = preprocessing.LabelEncoder()
        4     data_clean[col]= label.fit_transform(df_new[col])
```

```
In [24]: 1 data_clean.head(5)
```

```
Out[24]:
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeNumber	EnvironmentSatisfaction	...	Perform
0	41	1	2	1102	2	1	2	1	1	2	...	
1	49	0	1	279	1	8	1	1	2	3	...	
2	37	1	2	1373	1	2	2	4	4	4	...	
3	33	0	1	1392	1	3	4	1	5	4	...	
4	27	0	2	591	1	2	1	3	7	1	...	

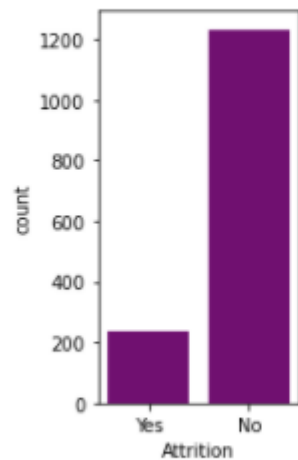
5 rows × 32 columns

HANDLING CLASS IMBALANCE

Classification problem where the distribution of examples across the known classes is biased or skewed. To avoid this we are using SMOTE technique

```
In [112]: 1 plt.figure(figsize=(2,4))
          2 sns.countplot(df.Attrition,color="purple")
```

```
Out[112]: <AxesSubplot:xlabel='Attrition', ylabel='count'>
```



the target column attrition has two values 0 and 1. It has class imbalance

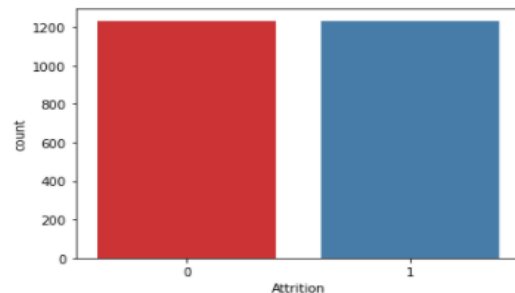
SMOTE synthetic over-sampling works to cause the classifier to build larger decision regions that contain nearby minority class points. This will in turn avoid data loss

```
In [333]: 1 x1=data.drop('Attrition',axis=1)
          2 y1=data['Attrition']
```

```
In [334]: 1 x1,y1=over.fit_resample(x1,y1)
```

```
In [335]: 1 sns.countplot(y1,palette="Set1")
```

```
Out[335]: <AxesSubplot:xlabel='Attrition', ylabel='count'>
```



Scaling using min max scaler

```
In [336]: 1 from sklearn.preprocessing import MinMaxScaler
          2 scaler=MinMaxScaler()
          3 scaled = scaler.fit_transform(x1)
```

Modelling

It is a binary classification problem so I have modelled using logistic regression and other classification models

__MODELING

```
1 from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
2 from sklearn.linear_model import LogisticRegression
3 from sklearn.ensemble import RandomForestClassifier
4 from sklearn.ensemble import BaggingClassifier
5 from sklearn.model_selection import train_test_split
6 from sklearn.neighbors import KNeighborsClassifier
7 from sklearn.metrics import roc_curve
8 from sklearn.metrics import roc_auc_score
9 from sklearn.model_selection import cross_val_score
10 from matplotlib import pyplot
11 from sklearn.svm import SVC
12 from sklearn.ensemble import GradientBoostingClassifier
13 from sklearn.tree import DecisionTreeClassifier
14 x=scaled
15 y=y1
```

```
1 train_test_split(x,y,test_size=0.25,random_state=1)
2
3 ),RandomForestClassifier(),BaggingClassifier(),KNeighborsClassifier(),GradientBoostingClassifier(),DecisionTreeClassifier()]
4
5
```

KNeighborsClassifier()

Accuracy score: 0.8

"Confusion Matrix:

" [[196 101]

[21 299]]

classification_report

	precision	recall	f1-score	support
0	0.90	0.66	0.76	297
1	0.75	0.93	0.83	320

accuracy			0.80	617
macro avg	0.83	0.80	0.80	617
weighted avg	0.82	0.80	0.80	617

Average accuracy_score 0.8022690437601296

BaggingClassifier()

Accuracy score: 0.84

"Confusion Matrix:

" [[262 35]

[63 257]]

classification_report

	precision	recall	f1-score	support
0	0.81	0.88	0.84	297
1	0.88	0.80	0.84	320

accuracy			0.84	617
macro avg	0.84	0.84	0.84	617
weighted avg	0.84	0.84	0.84	617

Average accuracy_score 0.8411669367909238

```
RandomForestClassifier()
```

```
Accuracy score: 0.9
```

```
"Confusion Matrix:
```

```
" [[276 21]
```

```
 [ 43 277]]
```

```
classification_report
```

	precision	recall	f1-score	support
0	0.87	0.93	0.90	297
1	0.93	0.87	0.90	320
accuracy			0.90	617
macro avg	0.90	0.90	0.90	617
weighted avg	0.90	0.90	0.90	617

```
Average accuracy_score 0.8962722852512156
```

```
LogisticRegression()
```

```
Accuracy score: 0.8
```

```
"Confusion Matrix:
```

```
" [[245 52]
```

```
 [ 69 251]]
```

```
classification_report
```

	precision	recall	f1-score	support
0	0.78	0.82	0.80	297
1	0.83	0.78	0.81	320
accuracy			0.80	617
macro avg	0.80	0.80	0.80	617
weighted avg	0.81	0.80	0.80	617

```
Average accuracy_score 0.8038897893030794
```

```
GradientBoostingClassifier()
```

```
Accuracy score: 0.88
```

```
"Confusion Matrix:
```

```
" [[268 29]
```

```
 [ 46 274]]
```

```
classification_report
```

	precision	recall	f1-score	support
0	0.85	0.90	0.88	297
1	0.90	0.86	0.88	320
accuracy			0.88	617
macro avg	0.88	0.88	0.88	617
weighted avg	0.88	0.88	0.88	617

```
Average accuracy_score 0.8784440842787682
```

```
DecisionTreeClassifier()

Accuracy score: 0.79

"Confusion Matrix:
" [[231  66]
  [ 64 256]]
classification_report
      precision    recall  f1-score   support

      0       0.78      0.78      0.78        297
      1       0.80      0.80      0.80        320

   accuracy          0.79          0.79          0.79          617
  macro avg          0.79          0.79          0.79          617
 weighted avg          0.79          0.79          0.79          617

Average accuracy_score 0.7893030794165316
```

Random forest classifier has highest accuracy is **0.896272**

Cross Validation:

In order to avoid over fitting, Cross-validation is used to estimate the skill of a machine learning model on unseen data.

```
In [51]: 1 score1=[]
```

```
In [52]: 1 lr=LogisticRegression()
2 scores=cross_val_score(lr,x,y,cv=5)
3 score1.append(scores)
4 scores
```

```
Out[52]: array([0.64574899, 0.85395538, 0.83975659, 0.86206897, 0.84178499])
```

```
In [53]: 1 rf=RandomForestClassifier()
2 scores=cross_val_score(rf,x,y,cv=5)
3 score1.append(scores)
4 scores
```

```
Out[53]: array([0.73481781, 0.95537525, 0.93103448, 0.95537525, 0.94523327])
```

```
In [54]: 1 bg=BaggingClassifier()
2 scores=cross_val_score(bg,x,y,cv=5)
3 score1.append(scores)
4 scores
```

```
Out[54]: array([0.70445344, 0.9127789 , 0.91075051, 0.90872211, 0.89655172])
```

```
In [55]: 1 kn=KNeighborsClassifier()
2 scores=cross_val_score(kn,x,y,cv=5)
3 score1.append(scores)
4 scores
```

```
Out[55]: array([0.76923077, 0.831643 , 0.81541582, 0.81541582, 0.82758621])
```

```
In [56]: 1 gb=GradientBoostingClassifier()
2 scores=cross_val_score(gb,x,y,cv=5)
3 score1.append(scores)
4 scores
```

```
Out[56]: array([0.65587045, 0.93711968, 0.90872211, 0.9148073 , 0.9148073 ])
```

Difference of predicted model and crossvalidation score:

- LogisticRegression() difference is 0.0378952
- RandomForestClassifier() difference is 0.05058173
- BaggingClassifier() difference is 0.06024702
- KNeighborsClassifier() difference is 0.02531716
- GradientBoostingClassifier() difference is 0.03636322
- DecisionTreeClassifier() difference is 0.05733428

from the observation KNeighborsClassifier model has least difference so I'm selecting KNeighborsClassifier as best model

Hyper Tuning:

__Hyper Tuning

```
In [226]: 1 from sklearn.model_selection import GridSearchCV, KFold
2 params = {
3     'n_neighbors' : [5,7,9,11,13,15],
4     'weights' : ['uniform','distance'],
5     'metric' : ['minkowski','euclidean','manhattan'],
6     'p':[1,2], 'leaf_size':list(range(1,20))
7 }
8
9
10 gs2 = GridSearchCV(KNeighborsClassifier(), params, verbose = 1, cv=3, n_jobs = -1)
11 gs2.fit(xtrain, ytrain)
12 print('Best param:', gs2.best_params_)

Fitting 3 folds for each of 1368 candidates, totalling 4104 fits
Best param: {'leaf_size': 1, 'metric': 'minkowski', 'n_neighbors': 5, 'p': 1, 'weights': 'distance'}
```

Best parameters: {'leaf_size': 1, 'metric': 'minkowski', 'n_neighbors': 5, 'p': 1, 'weights': 'distance'}

Modelling using best parameter and best model:

```

: 1 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=.25)
: 2 model = KNeighborsClassifier(metric='minkowski', n_neighbors=5, weights='distance', p=1, leaf_size=1)
: 3 model.fit(x_train, y_train)
: 4 model.score(x_test, y_test)

: 0.9027552674230146

: 1 y_pred_1 = model.predict(x_test)

: 1 result = confusion_matrix(y_test, y_pred_1)
: 2 print("Confusion Matrix:")
: 3 print(result)
: 4 result1 = classification_report(y_test, y_pred_1)
: 5 print("Classification Report:")
: 6 print(result1)
: 7 result2 = accuracy_score(y_test, y_pred_1)
: 8 print("Accuracy:", result2)

Confusion Matrix:
[[242  53]
 [  7 315]]
Classification Report:
              precision    recall  f1-score   support

     0       0.97       0.82       0.89       295
     1       0.86       0.98       0.91       322

 accuracy          0.91
 macro avg         0.91
 weighted avg      0.91

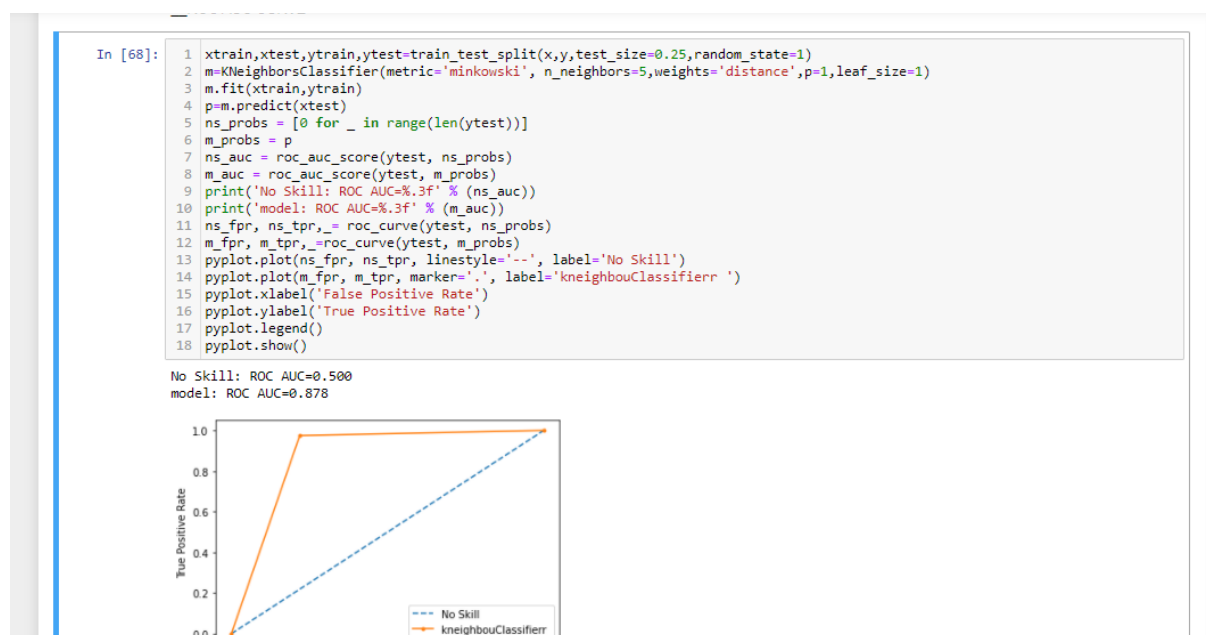
Accuracy: 0.9027552674230146

```

Final model after hyper tuning with accuracy **0.9027552674230146**

Best model: KNeighbourClassifier Best param: {'leaf_size': 1, 'metric': 'minkowski', 'n_neighbors': 5, 'p': 1, 'weights': 'distance'}

ROC AUC CURVE:



Conclusion:

I have developed a model to predict attrition of an employee with 90.2% accuracy

Saving the model

```
In [398]: 1 from joblib import dump  
          2 dump(model, 'model_hr.joblib')
```

```
Out[398]: ['model_hr.joblib']
```

```
In [399]: 1 from joblib import load  
          2 loaded = load('model_hr.joblib')
```